

# What is iptables?

IPTables is the default command-line-based stateless firewall for many Linux machines. In short, it is the thing that decides what happens to all incoming and outgoing packets. When a connection tries to establish itself to a system, IPTables matches it to a rule in its list(if none match default action is taken). Important note here is that the order of the rules in IPTables matters because it starts checking from the first rule and goes down the chain.

General IPTables format :

iptables → Tables → Chains → Rules

*iptables [-t table] command [match pattern] [action]*

\*IMPORTANT : if you are using ufw, its settings will override any iptables rules!

iptables tables	
filter	Used for normal filtering of traffic based on rules defined by the user(accept, reject, etc). Most used table in the iptables firewall. It is helpful in carrying out normal day-to-day blocking and filtering.
nat	Used for Network Address Translation purposes. NAT is a method of mapping an IP address space into another by modifying network address information in the IP header of packets while they are in transit across a traffic routing device.
mangle	Rules in this table can be used to modify the packers based on the user given criteria. It can be used to modify the TTL, MSS value, TOS(traffic priority), etc. Used to alter QOS bits in the TCP header.
raw	Primarily used to add No Connection Tracking rules.
security	Used for Mandatory Access Control networking rules.

iptables chains		
<b>INPUT</b> Chain used to control the behavior of incoming connections Ex. a user attempting to SSH into your server: iptables will match the user's IP address and port to a rule in the input chain.	<b>FORWARD</b> Chain used for incoming connections that aren't actually being delivered locally. Primarily used on pass-through devices like routers. Ex. NAT, routing, etc.	<b>OUTPUT</b> Chain used for outgoing connections Ex. you attempt to ping google.com : iptables will check what rules are regarding ping and google.com.
<b>POSTROUTING(NAT)</b> Alter packets before routing; packet translation happens immediately after the packet comes to the system(and before routing). An example use case is translating the destination IP address of the packet to something that matches the routing on the local machine; used for destination NAT(DNAT).	<b>PREROUTING(NAT)</b> Alter packets after routing; packet translation happens when the packers are leaving the system. Used for source NAT(SNAT).	<b>OUTPUT(NAT)</b> NAT for locally generated packets on the firewall

iptables target values	
ACCEPT	Firewall will accept the packet. Example of pinging INPUT ACCEPT :

	<pre> C:\Users\geek&gt;ping 192.168.6.129  Pinging 192.168.6.129 with 32 bytes of data: Reply from 192.168.6.129: bytes=32 time&lt;1ms TTL=64 Reply from 192.168.6.129: bytes=32 time&lt;1ms TTL=64 Reply from 192.168.6.129: bytes=32 time&lt;1ms TTL=64 Reply from 192.168.6.129: bytes=32 time&lt;1ms TTL=64  Ping statistics for 192.168.6.129:     Packets: Sent = 4, Received = 4, Lost = 0 (0% loss),     Approximate round trip times in milli-seconds:         Minimum = 0ms, Maximum = 0ms, Average = 0ms </pre>
DROP	<p>Firewall will drop the packet. Drop/ignore the connection. Use this option if you don't want the sender to realize your system exists.</p> <pre> C:\Users\geek&gt;ping 192.168.6.129  Pinging 192.168.6.129 with 32 bytes of data: Request timed out. Request timed out. Request timed out. Request timed out.  Ping statistics for 192.168.6.129:     Packets: Sent = 4, Received = 0, Lost = 4 (100% loss), </pre>
QUEUE	<p>Firewall will pass the packet to the userspace.</p>
RETURN / REJECT	<p>Firewall will stop executing the next set of rules in the current chain for this packet. The control will be returned to the calling chain. Use this option if you don't want the sender to establish a connection but want them to know that your firewall blocked them.</p> <pre> C:\Users\geek&gt;ping 192.168.6.129  Pinging 192.168.6.129 with 32 bytes of data: Reply from 192.168.6.129: Destination port unreachable. Reply from 192.168.6.129: Destination port unreachable. Reply from 192.168.6.129: Destination port unreachable. Reply from 192.168.6.129: Destination port unreachable.  Ping statistics for 192.168.6.129:     Packets: Sent = 4, Received = 4, Lost = 0 (0% loss), </pre>
MASQUERADE	<p>The masquerade option is used exactly the same as SNAT. The difference is each time that the MASQUERADE target gets hit by a packet, it automatically checks for the IP address to use-unlike SNAT, where the target is matched using single configured IP addresses. This means that the MASQUERADE option works properly with Dynamically allocated IP addresses,(DHCP).</p>

## Policy Chain DEFAULT behavior

When no matches are found for a specific packet, iptables will default to a set of pre-defined rules. To check what your iptables default policy chains are :

```
# iptables -L | grep policy
```

```
(root@pwner3000)-[/home/deusxmachina]
# iptables -L | grep policy
Chain INPUT (policy ACCEPT)
Chain FORWARD (policy ACCEPT)
Chain OUTPUT (policy ACCEPT)
```

To change the default policy chain rules :

```
# iptables --policy <chain-name> <target-value>
# iptables --policy INPUT DROP
# iptables --policy OUTPUT ACCEPT
# iptables --policy FORWARD DROP
# iptables -L | grep policy

(root@pwner3000)-[/home/deusxmachina]
# iptables -L | grep policy
Chain INPUT (policy DROP)
Chain FORWARD (policy DROP)
Chain OUTPUT (policy ACCEPT)
```

## Allowing or blocking specific connections

Aside from the default policy chains, you can configure additional rules to allow or block specific connections. You can either append a rule :

```
# iptables -A <chain-name> <rule-specification> [options]

{block all connections from IP address 10.0.0.10}
# iptables --append INPUT --source 10.0.0.10 -j DROP
# iptables -A INPUT -s 10.0.0.10 -j DROP // same command as above


{allow all connections to IP addresses 10.0.0.1 ~ 10.0.0.255}
# iptables --append OUTPUT --destination 10.0.0.0/24 -j ACCEPT
# iptables -A OUTPUT -d 10.0.0.0/24 -j ACCEPT // same command as above
# iptables -A OUTPUT -d 10.0.0.0/255.255.255.0 -j ACCEPT // same command as above

{block SSH connections from 10.0.0.10 - assuming SSH is running on port #22}
# iptables -A INPUT -p tcp --dport 22 -s 10.0.0.10 -j DROP
# iptables -A INPUT -p tcp --dport ssh -s 10.0.0.10 -j DROP // same command as above

{block SSH connections from anywhere}
# iptables -A INPUT -p tcp --dport ssh -j DROP
```

## NAT table examples


As shown in the 'iptables tables' table, the NAT table is used for mapping an IP address space into another by modifying network address information in the IP header of packets while they are in transit across a traffic routing device. This means that the NAT table option is used to change both the source address and destination address of a rule-matching packet. The relationship between the three predefined chains-PREROUTING, OUTPUT, and POSTROUTING-are as following :



Before any further NAT configurations, we need to activate IP-forwarding in the kernel :

```
# echo 1 > /proc/sys/net/ipv4/ip_forward //enable IP-forwarding
# echo 0 > /proc/sys/net/ipv4/ip_forward //disable IP-forwarding
```

In the following scenario, interface eth0 is connected to the local net, while eth1 is connected to the internet :



Packets arriving from the local net need their source IP address changed such that it is equal to the router's address. We configure on R1 :

```
# iptables -t nat -A POSTROUTING -o eth1 -j MASQUERADE
-t : select NAT table for configuration of NAT rules
-A POSTROUTING : append a rule to the POSTROUTING chain
-o eth1 : this rule is valid for packets that leave on the network interface 'eth1'
-j MASQUERADE : the action that should take place is to 'masquerade' packets; replace sender address with router's
```

## Choosing match patterns for NAT table

To match TCP packets **from** 10.0.0.1 :

```
# iptables -t nat -A POSTROUTING -p tcp -s 10.0.0.1 [action]
```

To match UDP packets going **to** 20.0.0.1 :

```
# iptables -t nat -A POSTROUTING -p udp -d 20.0.0.1 [action]
```

To match all packets from 30.0.0.0/16 arriving at eth0 :

```
# iptables -t nat -A PREROUTING -s 30.0.0.0/16 -i eth0 [action]
```

**IMPORTANT!** Remember that

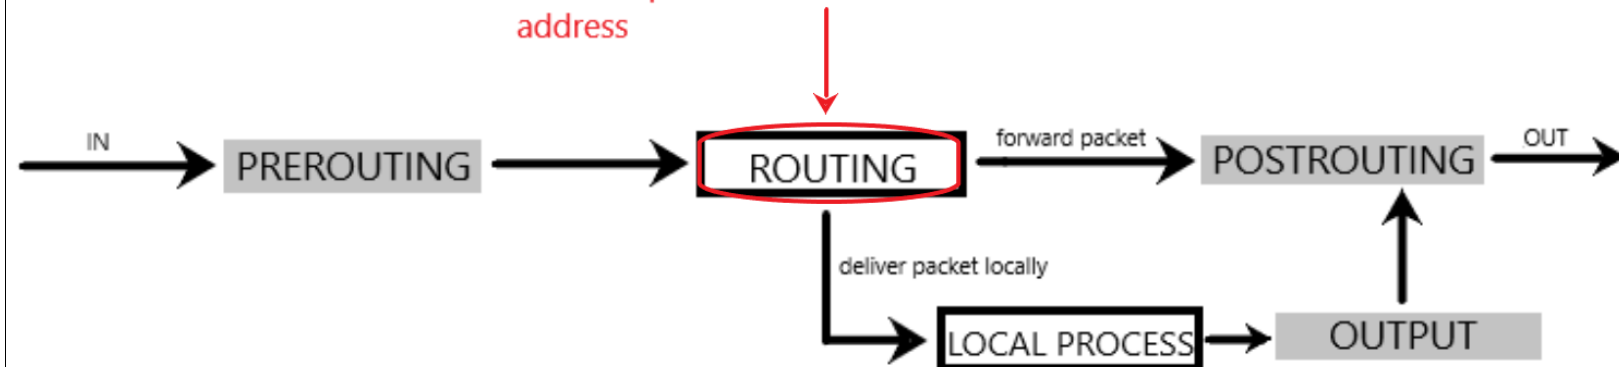
PREROUTING is used for DNAT for incoming packets

OUTPUT is used for DNAT for outgoing packets


POSTROUTING is used for SNAT for outgoing (local or forwarded) packets

An easy way to understand and memorize these distinctions is :


Routing by definition is choosing a route to deliver a packet based on it's DESTINATION address



So if you wanted to change the destination address of a packet you would have to do it before a routing decision based on the destination address can be done



And the changing the source address can be done once the route to the destination has been set



To match all packets except TCP packets and except packets from 192.168.1.2 :

```
# iptables -t nat -A PREROUTING -p ! tcp -s ! 192.168.1.2 [action]
```

To match packets leaving at eth1 :

```
# iptables -t nat -A POSTROUTING -o eth1 [action]
```

To match TCP packets from 192.168.1.2 ports 12345~12356 to 1.1.1.1 port 22 :

```
# iptables -t nat -A PREROUTING -p tcp -s 192.168.1.2 --sport 12345:12356 -d 1.1.1.1 --dport 22 [action]
```

## Actions for NAT table

To change sender address to 123.123.123.123 :

```
# iptables [match pattern] -j SNAT --to-source 123.123.123.123
```

To change sender address to outgoing network interface :

```
# iptables [match pattern] -j MASQUERADE
```

To change destination address to 123.123.123.123 port 22 :

```
# iptables [match pattern] -j DNAT --to-destination 123.123.123.123:22
```

To redirect packet to local port 8080 :

```
# iptables [match pattern] -j REDIRECT --to-ports 8080
```

## Solve two-way communication problems with 'Connection States'

You may have recognized a problem with the previous SSH filtering rules :

Many protocols require a two-way communication to function. Ex. when you ping [www.google.com](http://www.google.com) you need an OUTPUT rule allowing ping to the target AND an INPUT rule to retrieve the ping response. In the case of SSH, our previous OUTPUT and INPUT ACCEPT configuration would allow both external users to connect to our SSH server and us to connect to their SSH server. What if we want only external users to connect to our SSH (and block us from establishing a connection to their SSH)?

In the following example, we allow SSH connections FROM 10.0.0.1

```
# iptables -A INPUT -p tcp --dport ssh -s 10.0.0.1 -m state --state NEW,ESTABLISHED -j ACCEPT
```

But new SSH connections TO 10.0.0.1 are not. However, the system is permitted to send back information over SSH as long as the session has already been established (as in not initiating the connection).

```
# iptables -A OUTPUT -p tcp --sport 22 -d 10.0.0.1 -m state --state ESTABLISHED -j ACCEPT  
(default deny any)
```

## Save iptables configuration

To prevent current iptables configured rules from disappearing upon the next service reboot, run the following command (for Ubuntu)

```
# iptables-save
```

## Resources

[https://www.karlrupp.net/en/computer/nat\\_tutorial](https://www.karlrupp.net/en/computer/nat_tutorial)